

Aisa khanvan

Senior AI/ML Engineer | San Francisco, CA

Professional Summary

AI/ML Engineer with 5 years of experience building production NLP and LLM systems. Comfortable across the full stack from data pipelines to API deployment. Proven track record in developing Retrieval-Augmented Generation (RAG) architectures and deploying scalable transformer-based models in cloud environments.

Core Technical Skills

Category	Technologies
Languages	Python, SQL, Bash
AI & Machine Learning	Deep Learning, NLP, Transformers, LLMs, RAG
Frameworks	PyTorch, TensorFlow, scikit-learn, HuggingFace, LangChain
Data Science	pandas, numpy, scipy
Engineering & Deployment	FastAPI, REST APIs, Flask, Docker, Kubernetes, CI/CD, GitHub Actions
Infrastructure & DBs	AWS, PostgreSQL, Redis, Pinecone, FAISS
Version Control & Agile	Git, Agile Methodologies

Professional Experience

TechNova AI Solutions

Senior AI Engineer | 14 May 2026 – Present

- Architected and built a high-performance RAG-based customer support assistant using LangChain, OpenAI embeddings, and Pinecone, resulting in a 40% reduction in initial ticket response times.
- Engineered scalable ETL pipelines feeding feature stores for downstream ML models, ensuring data consistency across training and inference environments.

- Led the integration of GitHub Actions for automated CI/CD pipelines, streamlining the deployment of ML microservices to AWS.

DataStream Systems

Machine Learning Engineer | 7 May 2026 – 14 May 2026

- Successfully deployed transformer-based text classification models using FastAPI and Docker on AWS ECS, handling over 1 million requests daily.
- Optimized existing deep learning models using PyTorch, achieving a 15% improvement in inference latency.
- Collaborated with cross-functional teams in an Agile environment to deliver production-ready NLP features for enterprise clients.

Education

State University of Technology

Master of Science in Computer Science (Specialization in AI) | 21 May 2026

Global Institute of Engineering

Bachelor of Science in Software Engineering | 7 May 2026

Projects

Open Source LLM Wrapper

- Developed a lightweight Python library for interfacing with various LLM providers, currently hosting 500+ stars on GitHub.

Real-time Sentiment Engine

- Built a real-time sentiment analysis dashboard using Flask, Redis, and HuggingFace Transformers to monitor social media trends.